

### AlexNet – “The Big Dropout Net”

#### Rhyme:

“Two towers tall, ReLUs all;  
Big filters start, five-by-five call.  
Dropout guards when data’s small –  
AlexNet wins it all.”

#### What this helps you recall

- Two-stream GPU split → “two towers tall”
- ReLU activation → “ReLUs all”
- First conv uses large filters ( $11 \times 11$ , then  $5 \times 5$ ) → “big filters... five-by-five call”
- Dropout introduced → “dropout guards”
- ImageNet 2012 winner → “wins it all”

### ZFNet – “The Zoom-and-Fix Net”

#### Rhyme:

“Alex was wide, ZF zoomed inside;  
Seven-by-seven to see the stride.  
Deconv shows where features hide  
–  
Fix the net and ride.”

#### ✓ What this helps you recall

- Improved AlexNet → “Alex was wide... ZF zoomed inside”
- Changed  $11 \times 11 \rightarrow 7 \times 7$  first filter → “seven-by-seven”
- Reduced stride for better feature capture → “see the stride”
- Introduced deconv visualizations → “where features hide”

### GoogLeNet (Inception) – “The Multi-Path Net”

#### Rhyme:

“One net, many paths in sight;  
One-three-five in fusion tight.  
Bottlenecks make the model light –  
GoogLeNet keeps depth just right.”

#### ✓ What this helps you recall

- Inception module = many parallel paths → “many paths in sight”
- Uses  $1 \times 1$ ,  $3 \times 3$ ,  $5 \times 5$  kernels together → “one-three-five in fusion tight”
- Uses  $1 \times 1$  bottlenecks for dimensionality reduction → “bottlenecks... light”
- Very deep but efficient → “depth just right”

### ResNet – “The Skip-Connection Net”

#### Rhyme:

“If learning fails, skip the trail;  
Add it back so gradients sail.  
Hundreds deep without a bail –  
ResNet wins the trail.”

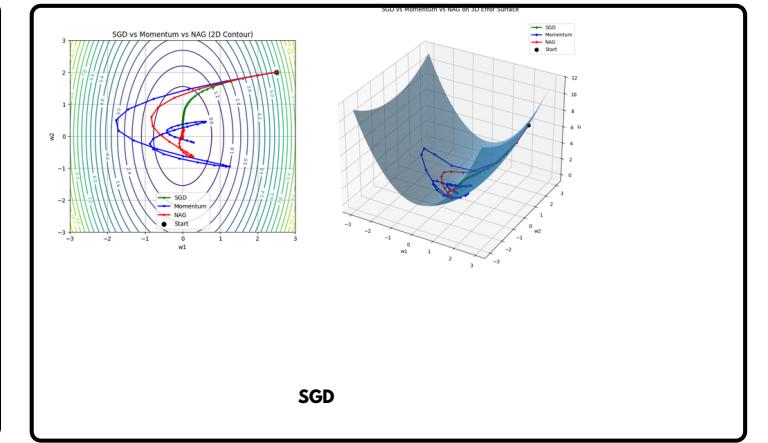
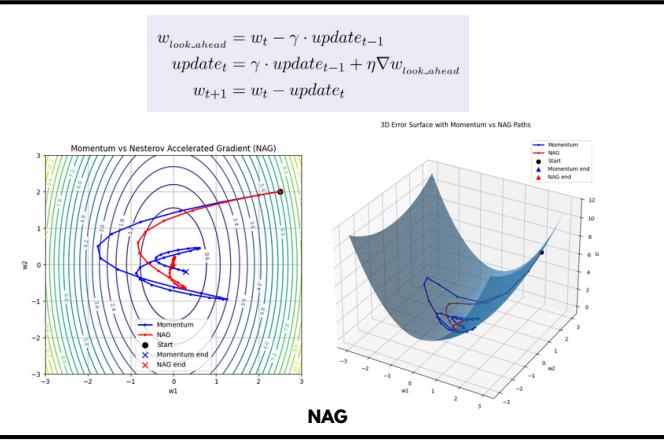
#### ✓ What this helps you recall

- Skip connections / identity mapping → “skip the trail”
- Solves vanishing gradient → “gradients sail”
- Extremely deep models ( $50/101/152$ ) → “hundreds deep”
- Breakthrough performance → “wins the trail”

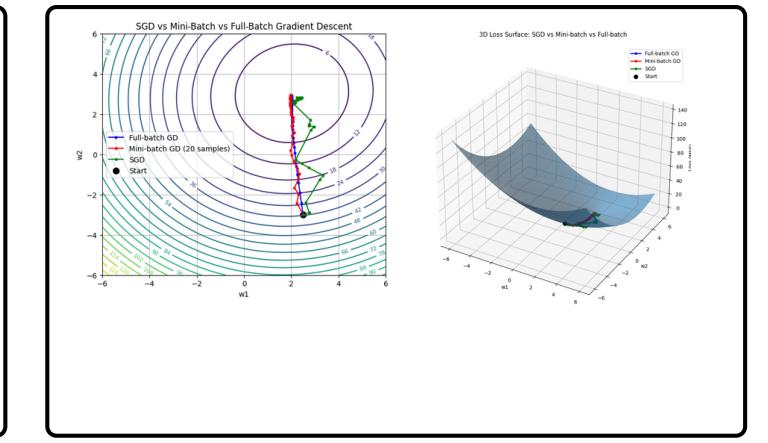
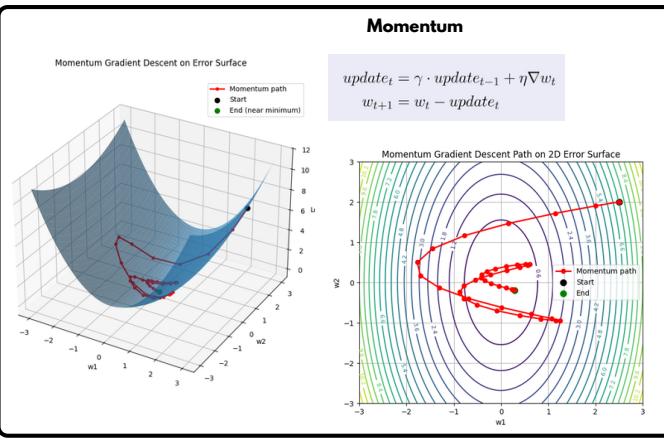
## Gradient Descend

“Batch is slow but steady in sight,  
SGD jumps left and right.  
Mini-batch makes learning tight,  
Momentum rolls with extra might,  
Nesterov looks ahead for the right-step  
light.”





Gradient descend



- Vanila Gradient Descend calculate gradient based on all sample points , then add all the calculated gradients and at last update the parameters like weight and biases using the summed gradient.
- The momentum and NAG variant of GD does the same only the update rule is different.
- In SGD, we calculate gradient on each sample point and update the parameters on each sample point. This does not lead to smooth convergence to optimal pint.
- The mini--batch version of SGD solved the problem by delaying the parameter update upto a predefined instance of gradient update

Gradient  $\longrightarrow g = \frac{1}{N} \sum_{i=1}^N \nabla \ell_i$

**Vanilla GD**      **Momentum**      **NAG**

**Look Ahead**

**NA**

**NA**

$$\tilde{w} = w - \eta \beta v$$

**Velocity**

**NA**

$$v = \beta v + g$$

$$v = \beta v + \nabla \ell(\tilde{w})$$

**Update**

$$w = w - \eta g$$

$$w = w - \eta v$$

$$w = w - \eta v$$