

# RNN



# Problem

● *Vanishing Gradient*

● *Exploding Gradient*

A Recurrent Neural Network (RNN) is a type of neural network designed to process sequential data by using feedback connections that allow information to persist across time steps, giving the network a form of memory.



# Solution

● *GRU*

● *LSTM*

## Vanishing Gradient

The vanishing gradient problem occurs when gradients become extremely small during backpropagation in deep or recurrent networks, causing earlier layers to learn very slowly or not at all.

- Happens mainly with sigmoid or tanh activations, where derivatives are small and repeatedly shrink during backpropagation.
- Deep networks and RNNs are most affected because gradients must pass through many layers or time steps.
- Leads to slow or stalled learning in early layers – network can't capture long-term dependencies.
- Causes models to favor short-term patterns, ignoring distant information.
- Solved or reduced using ReLU, LSTM/GRU gates, batch normalization, residual connections, and proper initialization.

## Exploding Gradient

The exploding gradient problem occurs when gradients become excessively large during backpropagation, causing unstable learning where model weights grow uncontrollably and training diverges.

- Most common in deep networks and RNNs where repeated multiplication of gradients causes them to blow up.
- Leads to numerical instability, causing loss values to become NaN or extremely large.
- Results in very large weight updates, making the model oscillate or fail to converge.
- Often caused by poor weight initialization or very long sequences in recurrent networks.
- Common solutions include gradient clipping, proper initialization (Xavier/He), using LSTM/GRU, and normalization techniques.

$$\frac{\partial L}{\partial h_t} = \left( \prod_{k=t+1}^T \frac{\partial h_k}{\partial h_{k-1}} \right) \frac{\partial L}{\partial h_T}$$

If each term

$$\left\| \frac{\partial h_k}{\partial h_{k-1}} \right\| = \alpha$$

then the gradient magnitude becomes:

$$\left\| \frac{\partial L}{\partial h_t} \right\| = \alpha^{(T-t)} \left\| \frac{\partial L}{\partial h_T} \right\|$$

- If  $\alpha < 1$ , then

$$\alpha^{(T-t)} \rightarrow 0 \Rightarrow \text{Vanishing gradient}$$

- If  $\alpha > 1$ , then

$$\alpha^{(T-t)} \rightarrow \infty \Rightarrow \text{Exploding gradient}$$

- $h_k$  is the hidden state of an RNN at time step  $k$ .
- It represents the memory or internal representation at that moment.

Example in RNN:

$$h_k = f(W_h h_{k-1} + W_x x_k)$$

- $L$  is the loss function (e.g., cross-entropy, MSE).
- During backpropagation, we compute how  $L$  changes with respect to each hidden state  $h_t$ .

Example:

$$L = \text{CrossEntropy}(y_{\text{true}}, y_{\text{pred}})$$

# A Comparative Analysis of GRU & LSTM

## GRU

A GRU (Gated Recurrent Unit) is a type of recurrent neural network that uses gates to control information flow, allowing it to learn long-term dependencies while avoiding vanishing gradients, similar to LSTM but with a simpler structure.

- Has two gates:
  - Update gate ( $z_t$ ) → decides how much of the past to keep
  - Reset gate ( $r_t$ ) → decides how much of the past to forget
- No separate cell state; the hidden state itself carries memory (simpler than LSTM).
- Faster and more efficient than LSTM because it has fewer parameters.
- Handles long-term dependencies better than simple RNNs.
- Works well for text, speech, time series, and any sequential task requiring memory.

$$z_t = \sigma(W_z x_t + U_z h_{t-1})$$

$$r_t = \sigma(W_r x_t + U_r h_{t-1})$$

$$\tilde{h}_t = \tanh(W_h x_t + U_h (r_t \odot h_{t-1}))$$

$$h_t = (1 - z_t) \odot h_{t-1} + z_t \odot \tilde{h}_t$$

## LSTM

A Long Short-Term Memory (LSTM) network is a type of recurrent neural network designed to learn long-term dependencies using a memory cell and three gates that control information flow, preventing vanishing and exploding gradients.

- Three gates:
  - Forget gate  $f_t$  → decides what to erase from memory
  - Input gate  $i_t$  → decides what new information to store
  - Output gate  $o_t$  → controls what part of memory becomes output
- Has a separate cell state  $C_t$  that carries long-term memory across many time steps.
- Designed to solve vanishing gradient problem, enabling learning of long sequences.
- More powerful but heavier and slower than GRU (more parameters).
- Used in NLP, speech recognition, time-series forecasting, and tasks requiring long-range context.

$$f_t = \sigma(W_f x_t + U_f h_{t-1})$$

$$i_t = \sigma(W_i x_t + U_i h_{t-1})$$

$$\tilde{C}_t = \tanh(W_C x_t + U_C h_{t-1})$$

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$

$$o_t = \sigma(W_o x_t + U_o h_{t-1})$$

$$h_t = o_t \odot \tanh(C_t)$$

